

# Towards Effective Attention Approximation in Biomedical Image Segmentation: Insights from ApproxDA-TransUNet

Author Name *Department of Computer Science*  
*University Name*  
 City, Country  
 author@university.edu

**Abstract**—Dual-attention mechanisms have demonstrated strong performance in medical image segmentation by jointly modeling long-range spatial and channel dependencies, but their quadratic complexity has motivated numerous approximation methods whose task-dependent behavior remains poorly understood. To bridge this gap, we propose Approximate Dual Attention TransUNet (ApproxDA-TransUNet), an efficient dual-attention approximation framework for hybrid CNN-Transformer architectures that replaces global attention with controllable windowed low-rank spatial attention and grouped channel attention. Unlike existing efficient attention methods that primarily pursue computational savings, ApproxDA-TransUNet also provides a controllable framework for systematically studying how attention approximation affects segmentation performance. Extensive experiments on Synapse, Kvasir-SEG, and ISIC 2018 demonstrate consistent improvements over DA-TransUNet, achieving 80.94%, 90.17%, and 89.58% Dice, respectively, but also reveal two critical insights: First, Global Context Sensitivity (GCS)—tasks requiring long-range dependency are far more sensitive to approximation scale than texture-dominated ones; Second, Gate Collapse—the commonly adopted learnable attention gate consistently degrades into near-uniform routing due to gradient symmetry, undermining adaptive fusion. These findings challenge the prevailing view of approximation as a pure efficiency tool, re-framing it as a task-dependent inductive bias that must be carefully calibrated to dataset characteristics. Our code and analysis offer both a strong segmentation baseline and actionable principles for designing efficient medical transformers.

**Index Terms**—medical image segmentation, dual attention, efficient transformers, attention approximation, global context sensitivity, gate collapse

## I. INTRODUCTION

Transformers have revolutionized medical image segmentation by enabling global long-range dependency modeling that convolutional networks cannot achieve. Dual attention [1]—combining a Position Attention Module (PAM) for spatial correlations and a Channel Attention Module (CAM) for channel correlations—has emerged as a powerful mechanism, and methods such as DA-TransUNet [2] demonstrated its effectiveness in hybrid CNN-Transformer architectures for multi-organ and polyp segmentation. The quadratic complexity of dual attention, however, has prompted a wave of efficient approximations—windowed attention [3], low-rank factorization [4], and grouped channel interaction [5]—that are now standard design choices. These strategies share an implicit

assumption that approximation is universally applicable with minor accuracy trade-offs—an assumption rarely examined across tasks with fundamentally different spatial reasoning requirements:

*When, and at what scale, is attention approximation beneficial for medical image segmentation?*

To investigate this, we design **ApproxDA-TransUNet**<sup>1</sup>: a dual-attention architecture with three orthogonal, independently controllable approximation operators—windowed PAM, low-rank projection, and grouped CAM. ApproxDA-TransUNet serves as both a practical efficient model that improves over DA-TransUNet on all four benchmarks, and an interpretable experimental platform for studying when and why approximation helps.

Applying ApproxDA-TransUNet across three medical segmentation benchmarks reveals a structured pattern: the same windowed attention operator is highly window-sensitive on multi-organ CT yet largely window-robust on polyp and skin lesion segmentation. Systematic ablation confirms that appropriate window selection consistently benefits all four benchmarks, but the degree of sensitivity differs markedly across tasks with different spatial context requirements. We further find that learnable routing collapses to a fixed equilibrium due to gradient symmetry—a structural property of symmetric two-branch architectures, not a training artifact.

The main contributions of this paper are:

- **ApproxDA-TransUNet**: We propose ApproxDA-TransUNet, an efficient dual-attention approximation architecture with three orthogonal controllable axes—windowed PAM, low-rank projection, and grouped CAM—that consistently improves segmentation performance over the original DA-TransUNet and competitive Transformer-based baselines.
- **Empirical discovery of GCS**: Using ApproxDA-TransUNet as a controllable experimental platform, we present a systematic empirical study of attention approximation across heterogeneous medical segmentation tasks, uncovering a  $3.6\times$  difference in window sensitivity between Synapse and Kvasir-SEG. We term this task-level variability *Global Context Sensitivity (GCS)* and

<sup>1</sup>Source code will be made available upon acceptance.

demonstrate that it predicts window-tuning importance across tasks.

- **Mechanistic insight:** We provide empirical evidence suggesting that attention approximation introduces a task-dependent inductive bias better aligned with tasks of lower contextual requirements, and further observe gradient symmetry collapse: learnable dual-branch routing degenerates to  $g \approx 0.5$  regardless of window size—a structural limitation of symmetric two-branch gating.

The pattern we uncover is structured by GCS: tasks with high contextual requirements demand careful window tuning; those with low requirements are largely window-robust. The underlying mechanism is an inductive bias shift—approximation does not just reduce cost, it changes what the attention mechanism is sensitive to.

The remainder of this paper is organized as follows: Section II reviews related work; Section III presents the ApproxDA-TransUNet framework; Section IV reports results against state-of-the-art methods; Section V provides ablation studies and mechanistic analysis; Section VI presents a generalization analysis of windowed attention as a task-dependent regularizer; Section VII investigates the GCS mechanism through causal elimination and introduces Semantic Spatial Dependency; Section VIII concludes.

## II. RELATED WORK

### A. CNN-based Medical Image Segmentation

Convolutional encoder-decoder architectures established the foundational paradigm for medical image segmentation, with U-Net [6] introducing the skip-connection design that underpins most segmentation networks, and successors [7]–[10] progressively adding dense connections, attention gating, and residual learning for improved feature reuse and training stability. These models offer strong computational efficiency and favorable inductive biases for local texture patterns, but their fixed receptive fields cannot capture long-range spatial dependencies across anatomical structures, limiting performance on tasks requiring global context.

### B. Transformer-based Medical Image Segmentation

Transformer architectures have significantly advanced medical image segmentation by enabling global dependency modeling beyond the locality of convolutional neural networks. Representative methods such as TransUNet [11], Swin-Unet [12], UNETR [13], UCTransNet [14], TransNorm [15], and DAE-former [16] combine convolutional feature extraction with transformer-based contextual modeling to achieve state-of-the-art segmentation performance across a wide range of medical imaging tasks. These methods treat global context modeling as uniformly beneficial, however, without examining whether different tasks require different degrees of spatial context.

### C. Dual Attention and DA-TransUNet

DANet [1] introduced Position Attention Modules (PAM) and Channel Attention Modules (CAM), capturing spatial and channel correlations at  $\mathcal{O}(N^2C)$  and  $\mathcal{O}(C^2N)$  complexity

respectively. DA-TransUNet [2] adapted this design for medical image segmentation by inserting dual-attention blocks at multiple skip connections of a hybrid R50-ViT-B/16 encoder-decoder, achieving state-of-the-art accuracy while inheriting the quadratic cost of full-map attention. DBAANet [17] further extends dual-branch attention by aggregating position and channel attention across multiple scales, demonstrating the continued appeal of dual-attention paradigms for medical image segmentation.

### D. Efficient Attention and Adaptive Routing

Windowed local attention [3], low-rank projection [4], and grouped channel interaction [5] have become standard strategies for reducing the cost of global attention, with learnable gating widely adopted to fuse multi-branch outputs [5], [16]. DAMamba [18] further shows that dual-attention routing transfers effectively to state-space models for lightweight breast ultrasound segmentation, confirming the versatility of multi-branch attention beyond standard transformer architectures. Existing efficient attention methods primarily optimize computational efficiency while implicitly assuming approximation is universally applicable. Whether approximation itself changes representation learning, and why it behaves differently across tasks with different contextual requirements, remains largely unexplored. This work addresses both questions by studying the context-sensitive behavior of approximation across heterogeneous segmentation tasks and exposing the gradient symmetry dynamics that cause adaptive routing to degenerate into fixed blending.

## III. METHOD

This section introduces the overall architecture of ApproxDA-TransUNet and the placement of ApproxDABlocks within it, then details the three independently controllable approximation axes each block implements, and finally analyzes their computational complexity.

### A. Architecture Overview

Figure 1 illustrates the overall architecture: ApproxDA-TransUNet adopts the hybrid CNN-ViT encoder-decoder structure of DA-TransUNet [2], with an R50-ViT-B/16 backbone pretrained on ImageNet-21K. Three CNN encoder stages produce multi-scale features at strides 4, 8, and 16; a 12-layer ViT processes  $14 \times 14$  patch tokens. ApproxDABlocks replace all dual-attention modules at three skip connections—512 channels at  $28 \times 28$ , 256 at  $56 \times 56$ , and 64 at  $112 \times 112$ —and at the bottleneck. Each block routes activations through a Low-Rank Windowed PAM branch and a Grouped CAM branch, fused by a configurable gate  $g$ . Figure 2 details the internal structure of each ApproxDABlock and its three independently controllable approximation axes.

### B. Attention Approximation Formulation

The original DA module performs dense spatial and channel attention over the full feature map, with quadratic complexity




Fig. 1. ApproxDA-TransUNet architecture. ApproxDABlocks (Low-Rank Windowed PAM + Grouped CAM + gate routing) replace dual-attention modules at all skip connections and the bottleneck. `gate_mode` controls the routing during ablation.




Fig. 2. ApproxDABlock internals with three independently controllable approximation axes: spatial scope ( $M$ , window size), representation fidelity ( $r$ , projection rank), and channel interaction ( $G$ , group count). Input  $x$  splits into a Low-Rank Windowed PAM branch (lavender, upper) and a Grouped CAM branch (green, lower); gate  $g$  fuses the outputs and  $x$  is added as a residual.

in both spatial resolution and channel dimension. ApproxDA-TransUNet reduces this cost along three independent axes, each targeting a distinct aspect of the approximation:

**Spatial scope** (window size  $M$ ) controls *how far* PAM reaches. Full global attention ( $M=H$ ) captures all long-range spatial relationships; restricting to  $M \times M$  local windows trades long-range context for lower complexity. This axis most directly governs whether tasks with broader contextual requirements retain the anatomical context they depend on.

**Representation fidelity** (rank  $r$ ) controls *how precisely* attention is computed within each window. Low-rank projection ( $r \ll M^2$ ) approximates the affinity matrix at reduced cost; higher  $r$  approaches exact attention.

**Channel interaction** (groups  $G$ ) controls *which channels* may interact in CAM. Partitioning  $C$  channels into  $G$  independent groups limits cross-channel dependency to within-group interactions, reducing the interaction scope.

These axes correspond to three distinct types of approx-

imation: *spatial approximation* (how far attention reaches, controlled by  $M$ ), *representation approximation* (how precisely affinity is computed, controlled by  $r$ ), and *channel approximation* (how broadly channels interact, controlled by  $G$ ). The axes are fully independent: any two of  $\{M, r, G\}$  can be held fixed while the third varies, enabling clean ablations that isolate individual approximation effects.

### C. Spatial Approximation

Instead of computing global position attention, ApproxDA-TransUNet partitions the feature map into non-overlapping windows of size  $M \times M$ . Attention is computed independently inside each local window, reducing complexity from  $\mathcal{O}(N^2C)$  to  $\mathcal{O}(NM^2C)$ .

**Window partitioning.** Following [3], we partition the feature map into  $n_W$  non-overlapping  $M \times M$  windows. Window clamping  $M \leftarrow \min(M, H, W)$  allows  $M=112$  to yield global attention at all scales, enabling a clean ablation that

isolates windowing from low-rank as the performance bottleneck.

**Low-rank projection.** Following [4], query and key features inside each window are projected from  $N_w=M^2$  to rank  $r \ll M^2$  via a learned matrix  $\mathbf{W}_r \in \mathbb{R}^{r \times N_w}$ :

$$\tilde{\mathbf{C}} = \mathbf{W}_r \mathbf{C}^\top, \quad \tilde{\mathbf{D}} = \mathbf{W}_r \mathbf{D}^\top \in \mathbb{R}^{r \times C}. \quad (1)$$

The attention matrix  $QK^\top \in \mathbb{R}^{N_w \times N_w}$  is approximated by  $Q_r K_r^\top \in \mathbb{R}^{N_w \times r}$ , where  $Q_r, K_r \in \mathbb{R}^{N_w \times r}$ , reducing total complexity to  $\mathcal{O}(NrC)$ . Combined, windowing and low-rank projection reduce the theoretical attention-matrix size from  $\mathcal{O}(N^2)$  to  $\mathcal{O}(Nr)$ , a factor of  $N/r$  ( $\sim 390\times$  at  $112 \times 112$  with  $r=32$ ). These two hyperparameters—window size  $M$  and rank  $r$ —form the primary experimental axes investigated on Synapse.

#### D. Channel Approximation

The original Channel Attention Module computes a  $C \times C$  channel attention matrix at  $\mathcal{O}(C^2N)$  cost. ApproxDA-TransUNet approximates this by partitioning  $C$  channels into  $G$  groups of width  $C_g = C/G$ , computing independent per-group attention:

$$X_{ji}^{(g)} = \frac{\exp(\mathbf{A}_i^{(g)} \cdot \mathbf{A}_j^{(g)})}{\sum_{k=1}^{C_g} \exp(\mathbf{A}_k^{(g)} \cdot \mathbf{A}_j^{(g)})}, \quad (2)$$

with total cost  $\mathcal{O}(C^2N/G)$ . Compared with dense channel interaction, grouped attention significantly reduces computational complexity while preserving local channel dependency.

The grouping number  $G$  provides another controllable approximation dimension, allowing the influence of channel approximation to be studied independently from spatial approximation.

#### E. Routing Strategy

The learnable routing mechanism of DA-TransUNet [2] is retained and extended to a configurable `gate_mode`. The outputs of the spatial and channel attention branches are fused via:

$$\mathbf{y} = \mathbf{W}_f \left( g \cdot \hat{\mathbf{P}} + (1 - g) \cdot \hat{\mathbf{C}} \right) + \mathbf{x}, \quad (3)$$

where  $\hat{\mathbf{P}}$  and  $\hat{\mathbf{C}}$  are the normalized PAM and CAM outputs,  $\mathbf{W}_f$  is a  $1 \times 1$  convolution, and  $g$  is determined by the `gate_mode` parameter:

$$\begin{aligned} \text{pam:} & \quad g = 1.0 \quad (\text{PAM only}) \\ \text{cam:} & \quad g = 0.0 \quad (\text{CAM only}) \\ \text{fixed:} & \quad g = 0.5 \quad (\text{equal blend}) \\ \text{learn:} & \quad g = \sigma(\mathbf{W}_g \cdot \text{GAP}(\mathbf{x})) \in (0, 1)^C \end{aligned}$$

In `gate=learn`,  $g \in (0, 1)^C$  is a per-channel routing coefficient broadcast over all spatial positions; it collapses to a channel-wise blend rather than a spatially-adaptive or input-dependent scalar. This single-parameter design enables systematic ablation of routing strategy without any other architectural change. ApproxDA-TransUNet does not treat this routing module as a primary architectural contribution; rather, it serves

as an analyzable component whose optimization dynamics—including an unexpected convergence to  $g \approx 0.5$  regardless of window size—are investigated in Section V.

The network is trained with a combined Dice and cross-entropy objective with equal weights:

$$\mathcal{L} = \frac{1}{2} \mathcal{L}_{\text{Dice}} + \frac{1}{2} \mathcal{L}_{\text{CE}}, \quad (4)$$

following DA-TransUNet [2] for direct comparability. Dice loss optimizes the overlap metric used at evaluation and handles class imbalance naturally; cross-entropy provides stable per-pixel gradient signal throughout training.

## IV. RESULTS

We evaluate ApproxDA-TransUNet on four benchmarks spanning tasks with fundamentally different spatial context requirements, comparing against DA-TransUNet and other competitive baselines.

#### A. Experimental Setup

**Datasets.** We evaluate on four representative medical image segmentation benchmarks.

*Synapse Multi-Organ CT* [19]: 30 abdominal CT scans with an 18/12 train/test split covering 8 organ classes. Metrics: mean DSC and HD95.

*Kvasir-SEG* [20]: 1000 endoscopy polyp images (binary) with an 800/200 train/test split. Metrics: DSC and mIoU.

*CVC-ClinicDB* [21]: 612 colonoscopy polyp images (binary) with a 489/123 train/test split. Metrics: DSC, mIoU, and HD95.

*ISIC 2018* [22]: 2594 dermoscopy skin lesion images (binary) with a 2075/519 train/test split. Metrics: DSC and mIoU.

**Implementation Details.** All experiments use an R50-ViT-B/16 backbone pretrained on ImageNet-21K, trained on NVIDIA Tesla T4 (16 GB) GPUs. Optimizer: SGD, lr=0.01; batch size 24; input  $224 \times 224$ ; 300 epochs; best checkpoint selected by evaluating every 15 epochs. Default ApproxDA-TransUNet configuration:  $M=7$ ,  $r=32$ ,  $G=8$ .

**Gate Protocol.** ApproxDA-TransUNet supports four gate modes: `pam` ( $g=1.0$ , PAM only), `cam` ( $g=0.0$ , CAM only), `fixed` ( $g=0.5$ , equal blend), and `learn` ( $g$  learned per channel). Unless otherwise specified, all main-result experiments use `gate=pam`.

#### B. Comparison with State-of-the-Art

Tables I–IV compare ApproxDA-TransUNet against published baselines across all four benchmarks.

**Synapse.** Table I benchmarks against 11 published methods. CNN-based models plateau at 76–78% DSC, with Att-UNet leading at 77.77%. Transformer-based methods progressively improve: TransUNet (77.48%), UCTransNet (78.23%), MIM (78.59%), Swin-Unet (79.13%), and DA-TransUNet (79.80%). ApproxDA-TransUNet with `gate=pam`,  $M=28$ ,  $r=32$  achieves **80.94%** DSC and 27.49 mm HD95—the highest DSC across all compared methods, surpassing DA-TransUNet by +1.14%. The intermediate window  $M=28$  captures sufficient cross-organ context while limiting spurious long-range correlations; Section V analyzes this behavior in detail.

TABLE I SEGMENTATION RESULTS ON SYNAPSE MULTI-ORGAN CT

| Method                                       | DSC (%) $\uparrow$ | HD95 (mm) $\downarrow$ |
|----------------------------------------------|--------------------|------------------------|
| <i>CNN-based</i>                             |                    |                        |
| U-Net [6]                                    | 76.85              | 39.70                  |
| U-Net++ [7]                                  | 76.91              | 36.93                  |
| Residual U-Net [9]                           | 76.95              | 38.44                  |
| MultiResUNet [10]                            | 77.42              | 36.84                  |
| Att-Unet [8]                                 | 77.77              | 36.02                  |
| <i>Transformer-based</i>                     |                    |                        |
| TransUNet [11]                               | 77.48              | 31.69                  |
| UCTransNet [14]                              | 78.23              | 26.75                  |
| TransNorm [15]                               | 78.40              | 30.25                  |
| MIM [23]                                     | 78.59              | 26.59                  |
| Swin-Unet [12]                               | 79.13              | <b>21.55</b>           |
| DA-TransUNet [2]                             | <u>79.80</u>       | <u>23.48</u>           |
| <i>Dual-attention approximation (ours)</i>   |                    |                        |
| <b>ApproxDA-TransUNet (ours)<sup>†</sup></b> | <b>80.94</b>       | 27.49                  |

<sup>†</sup> ApproxDA: gate=pam,  $M=28$ ,  $r=32$ ,  $G=8$ .

**Kvasir-SEG.** On polyp segmentation (Table II), U-Net leads the CNN-based group (88.22% DSC) while TransUNet (87.91%) and DA-TransUNet<sup>†</sup> (88.44%) are the strongest Transformer-based entries. ApproxDA-TransUNet with gate=pam,  $M=56$  achieves **90.17%** DSC, **84.27%** mIoU, and **44.35 mm** HD95—state-of-the-art across all three metrics, with a +1.73% DSC gain and −8.69 mm HD95 improvement over the nearest baseline.

TABLE II SEGMENTATION RESULTS ON KVASIR-SEG POLYP

| Method                           | DSC (%) $\uparrow$ | mIoU (%) $\uparrow$ | HD95 (mm) $\downarrow$ |
|----------------------------------|--------------------|---------------------|------------------------|
| U-Net [6]                        | 88.22              | 80.12               | —                      |
| Att-Unet [8]                     | 86.61              | 78.01               | —                      |
| U-Net++ [7]                      | 86.57              | 77.67               | —                      |
| Res-UNet [9]                     | 77.85              | 66.04               | —                      |
| TransUNet [11]                   | 87.91              | 80.03               | —                      |
| DA-TransUNet [2] <sup>†</sup>    | <u>88.44</u>       | <u>81.70</u>        | <u>53.04</u>           |
| <b>ApproxDA-TransUNet (ours)</b> | <b>90.17</b>       | <b>84.27</b>        | <b>44.35</b>           |

<sup>†</sup>We are able to closely reproduce paper-reported results (88.47% DSC, 81.02% mIoU) and capture the HD95. ApproxDA: gate=pam,  $M=56$ ,  $r=32$ .

**CVC-ClinicDB.** On colonoscopy polyp segmentation (Table III), DA-TransUNet<sup>‡</sup> achieves 89.47% DSC and 82.51% mIoU—the strongest published baseline on this dataset. ApproxDA-TransUNet with gate=pam,  $M=112$  achieves **90.99%** DSC, **85.09%** mIoU, and **14.77 mm** HD95—surpassing the nearest baseline by +1.52% DSC and +2.58% mIoU. The DSC span across  $M \in \{7, 28, 56, 112\}$  is 0.62 pp, confirming CVC-ClinicDB is window-robust—consistent with its binary task structure and low GCS, as analyzed in Section V.

**ISIC 2018.** On skin lesion segmentation (Table IV), CNN-based methods cluster between 83–88% DSC, with TransUNet (88.78%) and DA-TransUNet (88.88%) leading among Transformer-based approaches. ApproxDA-TransUNet with gate=learn,  $M=7$  achieves **89.58%** DSC (+0.70% over DA-TransUNet) and 82.66% mIoU, attaining the best DSC across

TABLE III SEGMENTATION RESULTS ON CVC-CLINICDB POLYP

| Method                           | DSC (%) $\uparrow$ | mIoU (%) $\uparrow$ | HD95 (mm) $\downarrow$ |
|----------------------------------|--------------------|---------------------|------------------------|
| U-Net [6]                        | 86.93              | 78.21               | —                      |
| Att-Unet [8]                     | 87.41              | 79.35               | —                      |
| U-Net++ [7]                      | 87.14              | 78.47               | —                      |
| Res-UNet [9]                     | 74.22              | 59.02               | —                      |
| TransUNet [11]                   | 89.01              | 81.63               | —                      |
| DA-TransUNet [2]                 | <u>89.47</u>       | <u>82.51</u>        | —                      |
| <b>ApproxDA-TransUNet (ours)</b> | <b>90.99</b>       | <b>85.09</b>        | <b>14.77</b>           |

ApproxDA: gate=pam,  $M=112$ ,  $r=32$ .

all compared methods. This window-robust pattern is consistent with the low-GCS behavior observed on Kvasir-SEG, analyzed in Section V.

TABLE IV SEGMENTATION RESULTS ON ISIC 2018 SKIN LESION

| Method                           | DSC (%) $\uparrow$ | mIoU (%) $\uparrow$ |
|----------------------------------|--------------------|---------------------|
| U-Net [6]                        | 87.22              | 81.14               |
| Att-Unet [8]                     | 87.60              | 81.51               |
| U-Net++ [7]                      | 87.30              | 81.33               |
| Res-UNet [9]                     | 83.32              | 76.51               |
| TransUNet [11]                   | 88.78              | 82.63               |
| DA-TransUNet [2]                 | <u>88.88</u>       | <u>82.78</u>        |
| <b>ApproxDA-TransUNet (ours)</b> | <b>89.58</b>       | <u>82.66</u>        |

ApproxDA: gate=learn,  $M=7$ ,  $r=32$ ,  $G=8$ .

ApproxDA-TransUNet achieves state-of-the-art results across all four benchmarks (Figure 3), with task-dependent window sensitivity emerging as the key differentiating pattern across datasets.




Fig. 3. Qualitative segmentation comparison (columns: input, ground truth, DA-TransUNet, ApproxDA). **Top (Synapse,  $M=28$ ):** multi-organ CT; ApproxDA 80.94% DSC (+1.14%). **Middle (Kvasir-SEG,  $M=56$ ):** polyp; 90.17% DSC (+1.73%). **Bottom (ISIC 2018,  $M=7$ ):** skin lesion; 89.58% DSC (+0.70%).

### C. Efficiency Analysis

Although the proposed approximation substantially reduces the theoretical complexity of the dual-attention module, the

TABLE V EFFICIENCY COMPARISON ON SYNAPSE (NVIDIA TESLA T4)

| Method                                  | Params (M) | GFLOPs | Train VRAM              | Train Time             | Infer (s/vol) | Multi-GPU       |
|-----------------------------------------|------------|--------|-------------------------|------------------------|---------------|-----------------|
| DA-TransUNet [2]                        | 107.95     | 30.2   | 11.5 G                  | 11.4h                  | 120.0         | No <sup>†</sup> |
| ApproxDA-TransUNet (gate=pam, $M=7$ )   | 112.98     | 32.1   | 6.4 G ( $\times 2$ GPU) | 8.3h ( $\times 2$ GPU) | 121.3         | Yes (DDP)       |
| ApproxDA-TransUNet (gate=learn, $M=7$ ) | 114.90     | 32.1   | 10.6 G                  | 11.5h                  | 114.2         | Yes (DDP)       |

<sup>†</sup>DA-TransUNet does not support multi-GPU training. ApproxDA-TransUNet uses Distributed Data Parallel (DDP); per-GPU batch 12, LR unchanged.

overall computational cost is dominated by the ViT backbone and convolutional encoder, which remain unchanged. As shown in Table V, end-to-end GFLOPs and memory exhibit only marginal differences from DA-TransUNet. The practical advantage of ApproxDA-TransUNet therefore lies in improved architectural compatibility rather than total model complexity reduction: DA-TransUNet fails under DDP, while ApproxDA-TransUNet trains natively across multiple GPUs, enabling cost-effective distributed training on commodity clusters.

## V. ABLATION STUDIES AND MECHANISTIC ANALYSIS

This section investigates why ApproxDA-TransUNet behaves differently under different approximation settings and segmentation tasks through a series of controlled ablation studies. Beyond validating the proposed design, these experiments provide mechanistic insights into attention approximation and its task-dependent behavior.

### A. Approximation Strength Ablation

We first investigate how approximation strength influences segmentation performance by varying the spatial window size and approximation rank while keeping all other components unchanged. Table VI isolates this effect by comparing  $M \in \{7, 14, 28, 56, 112\}$  under gate=pam,  $r=32$  on Synapse. Approximation is present in all rows—the right approximation *scale*, not the absence of approximation, governs performance.

TABLE VI SPATIAL APPROXIMATION ABLATION ON SYNAPSE ( $G=8$ )

| Config              | $M$ | $r$ | Gate               | DSC (%) $\uparrow$ | HD95 (mm) $\downarrow$ |
|---------------------|-----|-----|--------------------|--------------------|------------------------|
| Windowed + low-rank | 7   | 32  | pam                | 78.64              | 31.09                  |
| Windowed + low-rank | 14  | 64  | learn <sup>†</sup> | 78.04              | 29.09                  |
| Windowed + low-rank | 28  | 32  | pam                | <b>80.94</b>       | <u>27.49</u>           |
| Windowed + low-rank | 56  | 32  | pam                | 78.90              | 35.23                  |
| Global + low-rank   | 112 | 32  | pam                | <u>79.44</u>       | <b>27.26</b>           |

<sup>†</sup>Gate collapsed to  $g \approx 0.5$ .  $M=112$  via window clamping yields full-map attention at all scales.

The Synapse DSC-vs- $M$  curve is *non-monotonic*: peak at  $M=28$  (80.94%, +1.14% vs. DA-TransUNet), with all other scales underperforming—too-local ( $M=7$ : 78.64%), intermediate-upper ( $M=56$ : 78.90%), and fully-global ( $M=112$ : 79.44%). This pattern is analogous to the classical bias-variance tradeoff: under-context windows have high bias (missing necessary dependencies); over-context windows introduce spurious variance (irrelevant distant tokens); the intermediate  $M=28$  achieves the optimal balance. The DSC




Fig. 4.  $\Delta$ DSC vs. window size  $M$  (gate=pam,  $r=32$ ) for Synapse (red, high GCS) and Kvasir-SEG (green, low GCS). Annotation box: DSC span and  $3.6\times$  ratio.

range of 2.30 pp across  $M$  values confirms high window sensitivity on this task.

The  $M=14/r=64$  row shows that increasing rank does not compensate for gate collapse ( $g \approx 0.5$ , Section V-C): DSC 78.04% is lower than the gate=pam  $M=7/r=32$  baseline (78.64%), suggesting  $r=32$  already captures sufficient attention information and rank is not the performance bottleneck.

Preserving sufficient spatial receptive field is more important than increasing approximation rank; the optimal window scale—not approximation rank—governs segmentation performance.

### B. Cross-task Ablation and Global Context Sensitivity

We next examine whether the same approximation strategy behaves consistently across different segmentation tasks. Although ApproxDA-TransUNet uses identical approximation operators across all experiments, its sensitivity to window selection varies substantially. On Synapse multi-organ CT, the DSC-vs- $M$  curve (Figure 4) is non-monotonic with a sharp peak at  $M=28$  and a 2.30 pp range across the measured window sizes. On Kvasir-SEG, the curve is nearly flat (0.64 pp range) with a shallow peak at  $M=56$ , indicating the task is largely window-robust. On ISIC 2018, ApproxDA-TransUNet (gate=learn,  $M=7$ ) achieves 89.58% DSC vs. 88.88% for DA-TransUNet (+0.70%), consistent with the low-sensitivity pattern. With task-appropriate windows, gains are positive across all three benchmarks: +1.14% (Synapse,  $M=28$ ), +1.73% (Kvasir,  $M=56$ ), and +0.70% (ISIC,  $M=7$ ).

These observations suggest that segmentation tasks exhibit different degrees of **Global Context Sensitivity (GCS)**—the degree to which accurate segmentation depends on preserving large-scale spatial context. Since GCS is not directly



observable, we estimate it empirically through the performance response to window-size ablations, summarized by the DSC range:

$$\Delta\text{DSC} = \max_M \text{DSC}(M) - \min_M \text{DSC}(M),$$

measured under fixed  $\text{gate}=\text{pam}$ ,  $r=32$ . For Synapse,  $\Delta\text{DSC} = 2.30$  pp; for Kvasir-SEG, 0.64 pp—a  $3.6\times$  ratio confirming that Synapse exhibits substantially higher empirical GCS. Tasks with higher empirical GCS exhibit steep DSC-vs- $M$  curves with a well-defined optimal scale; lower-GCS tasks show broad performance plateaus where any reasonable window achieves competitive performance. Crucially, high GCS does not imply approximation is harmful: with the right window, ApproxDA-TransUNet outperforms full dual attention on *both* tasks. GCS characterizes how *sensitive* performance is to window choice—not the direction of the effect.

The task-level differences in GCS are grounded in the nature of each segmentation problem. Multi-organ abdominal CT segmentation requires reasoning over long-range anatomical relationships among multiple organs. The high sensitivity to window size reflects that both too-local windows (insufficient cross-organ context) and fully-global windows (noise from unrelated background) degrade performance; the intermediate scope ( $M=28$ ) provides the optimal balance. In contrast, polyp and skin lesion segmentation are primarily governed by local appearance, boundary contrast, and texture consistency. For these low-GCS tasks, attention approximation introduces an *inductive bias* favoring local interactions while suppressing unnecessary long-range dependencies—a prior that better matches the intrinsic structure of the problem.

We validate this locality bias hypothesis via controlled attention-map analysis: comparing  $\text{gate}=\text{pam}$   $M=7$  (windowed) versus  $M=112$  (global) attention effects on 10 Kvasir polyp images, windowed attention concentrates 62.3% of its effect on the polyp region versus only 34.2% for global attention—a  $1.82\times$  difference confirmed in 9/10 samples (Figure 5). Windowed attention focuses on polyp boundaries while global attention disperses to surrounding tissue, providing direct evidence for the inductive bias alignment hypothesis.

Together, these findings reframe the role of approximation: not as a capability reduction but as a task-dependent inductive bias adjustment whose value depends on how well the enforced locality prior aligns with the task’s intrinsic context scale.

### C. Gate Ablation and Mechanistic Analysis

Finally, we investigate whether adaptive routing contributes to approximation performance by comparing PAM-only, CAM-only, and learnable gate configurations. Table VII reports results for all four gate modes on Synapse ( $M=7$ ,  $r=32$ ).

Spatial attention alone provides stronger routing than either channel-dominated or learnable fusion.  $\text{Gate}=\text{pam}$  (78.64%) outperforms  $\text{gate}=\text{cam}$  (78.26%) and both  $\text{gate}=\text{learn}$  configurations (77.78% at  $M=7$ ; 78.04% at  $M=14$ ,  $r=64$ ). Notably,  $\text{gate}=\text{learn}$  with a larger window and higher rank ( $M=14$ ,  $r=64$ ) still underperforms  $\text{gate}=\text{pam}$  at  $M=7$  (78.04% vs. 78.64%),




Fig. 5. PAM contribution maps ( $\|\text{output} - \text{input}\|_2$ , per-panel normalised; colorbar below) on three Kvasir-SEG polyp images. Each row: input with GT contour (left), windowed PAM ( $M=7$ ; center), global PAM ( $M=112$ ; right). Brighter color indicates larger PAM contribution at that location; panels are independently normalised so magnitudes are not comparable across panels.

TABLE VII GATE MODE ABLATION ON SYNAPSE ( $M=7$ ,  $r=32$ )

| Gate Mode                                      | $g$             | DSC (%) $\uparrow$ | HD95 (mm) $\downarrow$ |
|------------------------------------------------|-----------------|--------------------|------------------------|
| $\text{gate}=\text{pam}$                       | 1.0             | <b>78.64</b>       | 31.09                  |
| $\text{gate}=\text{cam}$                       | 0.0             | 78.26              | 30.59                  |
| $\text{gate}=\text{learn}$ ( $M=14$ , $r=64$ ) | $\approx 0.5^*$ | 78.04              | <b>29.09</b>           |
| $\text{gate}=\text{learn}$                     | $\approx 0.5^*$ | 77.78              | 34.29                  |

\*Gate collapsed to  $g \approx 0.5$  throughout training.

confirming that gate collapse is more damaging than a sub-optimal window choice. In both  $\text{gate}=\text{learn}$  configurations, the gate converges to  $g \approx 0.5$  regardless of window size, producing a fixed 50/50 blend that performs below both single-branch modes.

The collapse is not isolated to one location but is uniform across the entire architecture. Figure 6 provides direct evidence across all four ApproxDABlock locations (top-left: bottleneck; top-right:  $28^2$  skip; bottom-left:  $56^2$  skip; bottom-right:  $112^2$  skip): gate values show no correlation with feature entropy (Spearman  $r_s \approx 0$ ), ruling out sample-adaptive routing. This collapse is identical at both  $M=7$  and  $M=14$ , confirming it is a fundamental property of symmetric two-branch gating rather than a windowing artifact.

**Optimization intuition.** Given fused output  $\mathbf{F} = g\mathbf{F}_P + (1 - g)\mathbf{F}_C$ , the gate gradient is:

$$\frac{\partial \mathcal{L}}{\partial g} = \frac{\partial \mathcal{L}}{\partial \mathbf{F}} \cdot (\mathbf{F}_P - \mathbf{F}_C). \quad (5)$$




Fig. 6. Gate value  $g$  vs. feature entropy  $H(F)$  at all four skip locations ( $M=7$ , gate=learn): bottleneck,  $28^2$ ,  $56^2$ ,  $112^2$ . All cluster at  $g \approx 0.5$  ( $\Delta g \leq 0.0002$ ), confirming gate collapse.

At initialization  $g \approx 0.5$ , both branches receive equal gradient  $\frac{1}{2} \partial \mathcal{L} / \partial \mathbf{F}$ . Without incentive to differentiate,  $\mathbf{F}_P \approx \mathbf{F}_C$ , so Eq. (5)  $\approx \mathbf{0}$ . This is self-reinforcing: equal gradients  $\rightarrow$  convergent representations  $\rightarrow$  near-zero gate gradient  $\rightarrow g$  fixed at 0.5—a stable equilibrium of symmetric dual-branch optimization.

**Empirical evidence.** Quantitative gate monitoring confirms this collapse is near-total across all window configurations. After 300 epochs, all four locations show  $\Delta g \leq 0.0002$  at both  $M=7$  and  $M=14$ , confirming the collapse is independent of window size.

Simple scalar gating cannot provide effective adaptive routing because branch symmetry prevents meaningful specialization. Future adaptive attention mechanisms should explicitly encourage branch diversity—through asymmetric initialization, diversity-encouraging loss terms, or non-symmetric routing—rather than relying solely on learnable fusion weights.

## VI. GENERALIZATION ANALYSIS: WINDOWED ATTENTION AS TASK-DEPENDENT REGULARIZER

Section V-B showed that ApproxDA-TransUNet consistently improves over DA-TransUNet across all three benchmarks with appropriate window selection. This section investigates a complementary question: does the locality prior introduced by windowed attention act as an implicit regularizer, and is its effect task-dependent?

### A. Optimization Volatility

We define *val DSC volatility* as the standard deviation of validation DSC measured at 15-epoch intervals over the final 30% of training (epochs 210–300, 7 evaluation points). Low volatility indicates stable convergence; high volatility indicates oscillatory optimization near the end of training. Table VIII reports volatility alongside best validation DSC, test DSC, and the val–test generalization gap for all six model–dataset combinations.

### B. CS-Dependent Volatility Crossover

The results reveal a consistent, task-dependent pattern (Figure 7). On low-GCS benchmarks, ApproxDA-TransUNet is dramatically more stable:  $18\times$  lower volatility on Kvasir-SEG (0.042 vs. 0.769) and  $9\times$  lower on ISIC 2018 (0.053 vs. 0.479). On high-GCS Synapse, this advantage *reverses*: ApproxDA-TransUNet is  $1.3\times$  more volatile than DA-TransUNet (0.542 vs. 0.409).

The sign reversal is the central finding. It is not explained by accuracy differences alone—ApproxDA achieves higher test DSC on all three tasks. The reversal suggests a mechanistic interpretation: windowed locality acts as a *stabilizing inductive prior* when it matches the task’s intrinsic context scale. On low-GCS tasks, local boundary cues suffice and windowed attention suppresses spurious long-range gradient signals, yielding smooth late-stage optimization. On high-GCS Synapse, the same locality constraint conflicts with the model’s need to aggregate cross-organ context, producing noisier gradients and higher volatility.

Importantly, neither model overfits in the conventional sense: all six val–test generalization gaps are  $\leq 0.72\%$ , and test DSC consistently meets or exceeds best validation DSC. The volatility difference is therefore a property of the optimization landscape, not a train–test distribution mismatch.

A comprehensive volatility sweep across all Synapse ApproxDA configurations confirms the pattern is robust: the gate=pam family at  $M \in \{7, 28, 112\}$  consistently shows  $1.1\text{--}1.3\times$  higher volatility than DA-TransUNet on Synapse, with gate=cam as the sole exception ( $0.44\times$ —consistent with channel-wise attention being less sensitive to spatial window size).

These results extend the GCS framing from a performance metric to an optimization phenomenon: locality prior alignment with task structure predicts both final accuracy (Section V-B) and training stability. Low-GCS tasks benefit doubly—from inductive bias alignment and from the implicit regularization effect of windowed attention.

## VII. GCS MECHANISM: CAUSAL ELIMINATION AND SEMANTIC SPATIAL DEPENDENCY

Section V-B characterized GCS empirically through window-size ablations. This section investigates the underlying mechanism through a systematic causal elimination study: five candidate explanatory factors are computed from ground-truth masks alone (no training required), assessed against the known GCS ordering, and ruled out one by one, converging on Semantic Spatial Dependency (SSD) as the unifying explanation.

### A. Candidate Metrics and Evaluation Protocol

We compute the following metrics from 200 randomly sampled training masks per dataset:

- **SC1 — Foreground Ratio:** Fraction of pixels belonging to any foreground class. Tests whether GCS is driven by target scale.
- **SC2 — Boundary Complexity:** Mean isoperimetric ratio ( $P^2/4\pi A$ ) per connected component. Tests whether GCS is driven by boundary irregularity.



TABLE VIII OPTIMIZATION VOLATILITY AND GENERALIZATION GAP (DSC, %). VOLATILITY = STD OF VAL DSC OVER THE LAST 30% OF TRAINING EPOCHS (EPOCHS 210–300); LOWER IS MORE STABLE.

| Dataset | GCS Level | Model               | Best Val (%) | Test DSC (%) | Val–Test Gap | Volatility   |
|---------|-----------|---------------------|--------------|--------------|--------------|--------------|
| Synapse | High      | DA-TransUNet        | 79.52        | 79.80        | −0.28        | 0.409        |
| Synapse | High      | ApproxDA ( $M=28$ ) | 81.02        | 80.94        | +0.08        | 0.542        |
| Kvasir  | Low       | DA-TransUNet        | 88.38        | 88.44        | −0.06        | 0.769        |
| Kvasir  | Low       | ApproxDA ( $M=7$ )  | 89.23        | 89.24        | −0.01        | <b>0.042</b> |
| ISIC    | Low       | DA-TransUNet        | 87.71        | 88.43        | −0.72        | 0.479        |
| ISIC    | Low       | ApproxDA ( $M=7$ )  | 89.56        | 89.58        | −0.02        | <b>0.053</b> |

TABLE IX CAUSAL ELIMINATION RESULTS ( $M=28$ ,  $n=200$  MASKS/DATASET, FIVE DATASETS INCLUDING ACDC DISCRIMINATING TEST).  $\rho$  = SPEARMAN RANK CORRELATION VS. KNOWN  $\Delta$ DSC ORDERING. A SUCCESSFUL RULE-OUT REQUIRES  $|\rho| \approx 0$  OR WRONG-DIRECTION (NEGATIVE) CORRELATION.

| Check | Metric               | Synapse | ACDC           | Kvasir      | ISIC        | CVC         | $\rho$             | Verdict                |
|-------|----------------------|---------|----------------|-------------|-------------|-------------|--------------------|------------------------|
| SC1   | Foreground ratio     | 0.077   | 0.040          | 0.152       | 0.214       | 0.094       | −0.50              | Ruled out              |
| SC2   | Boundary complexity  | 1.355   | 0.928          | 1.023       | 1.261       | 1.163       | +0.30              | Ruled out <sup>†</sup> |
| SC3   | Spatial entropy      | 0.497   | 0.359          | 0.566       | 0.605       | 0.482       | −0.10              | Ruled out              |
| SC4   | WCR ( $M=28$ px)     | 0.999   | 0.954          | 1.000       | 0.999       | 1.000       | −0.60              | Ruled out              |
| SC5   | Inter-class sep.     | 0.9996  | 0.547          | 0 (constr.) | 0 (constr.) | 0 (constr.) | +0.89              | Strongest predictor    |
| Proxy | $n_{\text{classes}}$ | 3.99    | $\approx 3.00$ | 1.00        | 1.00        | 1.00        | +0.87 <sup>‡</sup> | Best proxy (3-dataset) |

<sup>†</sup>SC2  $\rho = +0.30$  with five datasets; CT boundary irregularity (imaging physics) inflated the 3-dataset result to  $\rho = +1.00$  but does not survive multi-dataset inclusion. <sup>‡</sup> $n_{\text{classes}}$   $\rho$  computed on original three datasets; ACDC ( $n_{\text{classes}} \approx 3$ , GCS=0.73 pp) would further suppress this correlation, confirming class count is a proxy, not a cause.

- **SC3 — Spatial Entropy:** Normalized entropy of the foreground pixel distribution over an  $8 \times 8$  spatial grid. Tests whether GCS is driven by foreground dispersion.
- **SC4 — Window Crossing Ratio (WCR):** Fraction of foreground pixels belonging to components spanning more than one  $M \times M$  window ( $M=28$  pixels, the pixel-space equivalent of the model’s  $M=7$  feature-space window at  $4 \times$  Swin downsampling). Tests whether individual target components crossing window boundaries drives GCS.
- **SC5 — Inter-class Window Separation:** Fraction of class-centroid pairs falling in different  $M \times M$  windows (multi-class datasets; zero by construction for binary tasks). Measures whether distinct semantic structures require cross-window attention.

For SC1–SC4 (rule-out checks), a successful elimination requires Spearman  $\rho \approx 0$  or wrong-direction correlation against the known GCS ordering (Synapse 2.30 pp > ACDC 0.73 pp  $\approx$  ISIC 0.70 pp  $\approx$  Kvasir 0.64 pp  $\approx$  CVC 0.62 pp, five datasets). An additional proxy search identifies  $n_{\text{classes}}$ —the mean number of foreground semantic classes per image—as a training-free GCS predictor. Table IX summarizes all results.

### B. Ruling Out Alternative Explanations

**SC1 and SC3 are ruled out.** Foreground ratio ( $\rho = -0.50$ ) and spatial entropy ( $\rho = -0.10$ ) show near-zero or wrong-direction correlations across all five datasets. Synapse has the *smallest* foreground coverage (7.7%) and the *lowest* spatial entropy, yet the highest GCS. GCS is not driven by how large or how uniformly distributed the foreground is.

**SC4 saturation rules out component-level window crossing.** At  $M=28$  pixels, WCR saturates near 1.0 for all five datasets ( $\rho = -0.60$ , wrong direction). Even a single large polyp crosses window boundaries in pixel space, independent of task type. The saturation itself is informative: *whether an individual target region crosses window boundaries does not predict GCS*. The discriminative mechanism must operate at a higher semantic level than individual components.

**SC2 does not survive the five-dataset test.** On three datasets, boundary complexity ranked Synapse > ISIC > Kvasir ( $\rho = +1.00$ ). Adding ACDC and CVC drops this to  $\rho = +0.30$ , ruling it out. The three-dataset result was an imaging-modality confound: CT physics produces high boundary contrast regardless of class count, which happened to correlate with GCS when all three datasets were compared—ACDC and CVC break this spurious alignment.

### C. Semantic Spatial Dependency as the Unifying Mechanism

The surviving evidence points to a single mechanism. SC5—inter-class window separation—achieves  $\rho = +0.89$  across five datasets, the strongest predictor among all five metrics. Synapse’s SC5 is 0.9996: windowed attention at  $M=28$  cannot simultaneously attend to nearly all organ pairs. ACDC’s SC5 is 0.547: cardiac structures are partially separated at this scale. Binary datasets have SC5=0 by construction. The proxy metric  $n_{\text{classes}}$  ( $\rho = +0.87$ , three datasets) captures the same structure but is not the cause—ACDC ( $n_{\text{classes}} \approx 3$ , high as Synapse yet low GCS) directly falsifies it.

We term the underlying mechanism *Semantic Spatial Dependency (SSD)*: the extent to which accurate segmentation




Fig. 7. Validation DSC trajectories (left column) and training loss trajectories (right column) for DA-TransUNet vs. ApproxDA-TransUNet on Synapse (top, high GCS), Kvasir (middle, low GCS), and ISIC (bottom, low GCS). ApproxDA achieves 9–18 $\times$  lower volatility on low-GCS tasks and slightly higher volatility on high-GCS Synapse, consistent with the inductive bias interpretation.

of each class requires locating it relative to other semantic structures. Multi-organ CT exemplifies high SSD—the pancreas must be located relative to the stomach, duodenum, and adjacent vessels; no single window captures all cross-organ context simultaneously. Binary segmentation tasks have zero inter-class SSD by definition. The causal chain is:

$$\begin{aligned}
 \underbrace{\text{SSD}}_{\text{task property}} &\rightarrow \underbrace{\text{cross-structure reasoning}}_{\text{attention requirement}} \\
 &\rightarrow \underbrace{\text{sensitivity to window scale.}}_{= \text{GCS}}
 \end{aligned}$$

The empirical proxy  $n_{\text{classes}}$  correlates with SSD because

more classes generally implies richer inter-class spatial dependencies, but class count is not the cause—and this distinction is directly validated by an additional experiment.

**ACDC as a discriminating test case.** To separate  $n_{\text{classes}}$  from SSD as the true driver, we ran the same window ablation ( $M \in \{7, 28, 56, 112\}$ , gate=pam) on the ACDC cardiac MRI benchmark [24], a 4-class task (background, RV, Myo, LV). If class count drives GCS, ACDC should show intermediate-to-high sensitivity. Instead,  $\Delta\text{DSC}=0.73$  pp—indistinguishable from the binary datasets Kvasir-SEG (0.64 pp), CVC-ClinicDB (0.62 pp), and ISIC 2018 (0.70 pp), and far below Synapse (2.30 pp).

SC5 quantifies this directly. ACDC’s inter-class window




Fig. 8. Causal elimination results ( $M=28$ , 200 masks/dataset, five datasets). Top row: per-dataset metric values. Bottom row: scatter vs. known  $\Delta$ DSC (GCS). SC1 ( $\rho = -0.50$ ) and SC3 ( $\rho = -0.10$ ) show near-zero or wrong-direction correlations; SC4 saturates near 1.0 across all datasets ( $\rho = -0.60$ ); SC2’s confound ( $\rho = +1.00$  on three datasets) collapses to  $\rho = +0.30$  with five datasets. All four alternative explanations are eliminated; SC5 ( $\rho = +0.89$ ) is the strongest predictor of GCS.




Fig. 9. Mask-based proxy metrics vs. known  $\Delta$ DSC.  $n_{\text{classes}}$  (leftmost) shows the strongest Spearman correlation ( $\rho = +0.87$ ). Spatial metrics (centroid spread, bbox coverage) show wrong-direction correlations, consistent with the SC1/SC3 elimination.

separation is 0.547 at  $M=28$ —cardiac structures (RV, Myo, LV) occupy adjacent but overlapping windows due to their concentric anatomy—compared to Synapse’s near-complete separation ( $SC5=0.9996$ ) and binary tasks’ structural zero. Crucially, GCS exhibits a strongly non-linear response: the drop from  $SC5=0.9996$  to 0.547 (Synapse  $\rightarrow$  ACDC) produces a 1.57 pp fall in GCS (2.30  $\rightarrow$  0.73 pp), while the further drop from 0.547 to 0 (ACDC  $\rightarrow$  binary tasks: Kvasir 0.64 pp, CVC 0.62 pp, ISIC 0.70 pp) yields negligible additional change ( $\approx 0.09$ –0.11 pp). Near-complete inter-class separation—not merely multi-class structure—is what triggers high GCS.

The ACDC result therefore refines the SSD theory: the key

TABLE X EMPIRICAL GCS ( $\Delta$ DSC) ACROSS FIVE DATASETS. ACDC AND CVC-CLINICDB BOTH FALL IN THE LOW-GCS TIER DESPITE HAVING 4 CLASSES AND A DIFFERENT DATASET SCALE RESPECTIVELY, CONFIRMING THAT TASK STRUCTURE ( $SC5$ ), NOT CLASS COUNT OR DATASET SIZE, DETERMINES GCS TIER.

| Dataset      | Task                | Classes | $\Delta$ DSC (pp) |
|--------------|---------------------|---------|-------------------|
| Synapse      | Multi-organ CT      | 9       | 2.30              |
| ACDC         | Cardiac MRI         | 4       | 0.73              |
| ISIC 2018    | Skin lesion         | 2       | 0.70 <sup>†</sup> |
| Kvasir-SEG   | Polyp (endoscopy)   | 2       | 0.64              |
| CVC-ClinicDB | Polyp (colonoscopy) | 2       | 0.62              |

<sup>†</sup>ISIC 2018  $\Delta$ DSC pending full gate=pam window ablation; 0.70pp is a provisional lower bound ( $\Delta$  vs. DA-TransUNet at  $M=7$ ).

quantity is not class count per se but *degree of inter-class spatial separation relative to the attention window*. A multi-class task with compact, stereotyped anatomy ( $SC5 < 0.6$ ) behaves identically to a binary task in terms of GCS.  $SC5$ , computable from annotation masks alone, provides a training-free SSD estimate that predicts GCS tier without running a single experiment.

**Practical implication.** GCS-aware window selection reduces to an annotation-level check: tasks with near-complete inter-class window separation ( $SC5 > 0.8$ , e.g. multi-organ CT) are predicted high-GCS and require window tuning. Tasks with low or intermediate  $SC5$ —whether binary or multi-class with compact anatomy—are predicted low-GCS and largely window-robust.

**Per-organ validation within Synapse.** The global GCS ( $\Delta$ DSC=2.30pp) aggregates over eight organs whose individual window sensitivities span 0.39–6.45 pp (Table XI),

TABLE XI PER-ORGAN DSC (%) AND WINDOW SENSITIVITY ( $\Delta$ DSC) ON SYNAPSE (GATE=PAM,  $r=32$ ,  $n=4$  WINDOW SIZES). SORTED BY  $\Delta$ DSC DESCENDING. PANCREAS EXCEPTION NOTED IN TEXT.

| Organ       | $M=7$ | $M=28$ | $M=56$ | $M=112$ | $\Delta$ DSC |
|-------------|-------|--------|--------|---------|--------------|
| Gallbladder | 65.2  | 68.6   | 62.2   | 64.9    | 6.45 pp      |
| Stomach     | 73.8  | 78.7   | 75.5   | 77.0    | 4.97 pp      |
| Kidney (R)  | 78.1  | 82.2   | 82.1   | 80.5    | 4.16 pp      |
| Spleen      | 87.5  | 89.4   | 85.9   | 86.4    | 3.52 pp      |
| Pancreas    | 60.6  | 62.7   | 61.6   | 62.7    | 2.05 pp      |
| Kidney (L)  | 82.5  | 84.5   | 84.4   | 82.5    | 2.02 pp      |
| Aorta       | 87.6  | 88.0   | 86.1   | 87.8    | 1.89 pp      |
| Liver       | 93.7  | 93.3   | 93.4   | 93.6    | 0.39 pp      |
| Mean        | 78.6  | 80.9   | 78.9   | 79.4    | 2.30 pp      |

providing within-task evidence for SSD at the organ level. Gallbladder shows the highest sensitivity (6.45 pp): it is small, partially occluded under the liver, and its localization requires the global liver–gallbladder spatial relationship—a textbook high-SSD structure. Stomach (4.97 pp) has highly variable shape and fill state, requiring surrounding organ context for boundary delineation. By contrast, the liver (0.39 pp) is the dominant abdominal landmark—identifiable at any window scale without cross-organ cues—and the aorta (1.89 pp) is a cylindrical structure running along a fixed anatomical axis, similarly window-robust.

The pancreas is a special case: it achieves only 60–63% DSC across *all* four window sizes, so its  $\Delta$ DSC (2.05 pp) reflects a consistent approximation ceiling rather than low SSD. Where the model achieves reasonable baseline performance,  $\Delta$ DSC reliably tracks anatomical SSD.

### VIII. CONCLUSION

This paper presented ApproxDA-TransUNet, an efficient dual-attention approximation framework for medical image segmentation. Across three representative benchmarks, ApproxDA-TransUNet consistently outperformed the original DA-TransUNet while retaining efficient attention computation, demonstrating that appropriately designed attention approximation can improve segmentation accuracy rather than merely reducing computational cost. Beyond the proposed architecture, our experiments provide new insights into attention approximation. We show that approximation effectiveness depends on the contextual requirements of the target task, which we describe as Global Context Sensitivity (GCS). High-GCS tasks require careful approximation design, whereas low-GCS tasks remain robust across a broad range of approximation scales. We further demonstrate that conventional scalar gate learning collapses toward nearly uniform routing because of gradient symmetry, suggesting that effective adaptive fusion requires mechanisms that explicitly promote branch specialization. Beyond the core results, this journal extension presents two additional analyses. A training dynamics study (Section VI) shows that windowed attention acts as a task-dependent regularizer: on low-GCS tasks it reduces val DSC volatility by 9–18 $\times$  relative to global attention, while slightly increasing volatility on high-GCS Synapse. This CS-dependent stability crossover is consistent with the inductive

bias interpretation and holds across all gate=pam window configurations. A causal elimination study (Section VII) systematically rules out image texture, object size, spatial dispersion, and component-level window crossing as GCS drivers, converging on Semantic Spatial Dependency (SSD)—the requirement to locate each semantic structure relative to others—as the unifying mechanism. A discriminating experiment on ACDC cardiac MRI (4 classes, spatially co-located anatomy) yields  $\Delta$ DSC=0.73 pp, identical to binary tasks and far below Synapse (2.30 pp), directly confirming that spatial spread of classes—not class count—governs GCS.

Overall, attention approximation should be regarded as a task-dependent inductive bias, not solely an efficiency technique. The primary open direction is extending the GCS spectrum to tasks with intermediate spatial dependency structures (e.g., multi-class datasets with partially distributed anatomy), and designing attention routing mechanisms that explicitly promote branch diversity to avoid gradient symmetry collapse. Future research should first ask *how much context* a task actually requires—a question GCS and SSD can answer from annotation structure alone—before deciding how to approximate.

### REFERENCES

- [1] J. Fu *et al.*, “Dual attention network for scene segmentation,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 3146–3154.
- [2] G. Sun *et al.*, “DA-TransUNet: Integrating spatial and channel dual attention with transformer U-Net for medical image segmentation,” *Frontiers in Bioengineering and Biotechnology*, vol. 12, p. 1398237, 2024.
- [3] Z. Liu *et al.*, “Swin transformer: Hierarchical vision transformer using shifted windows,” in *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021, pp. 10012–10022.
- [4] S. Wang *et al.*, “Linformer: Self-attention with linear complexity,” *arXiv preprint arXiv:2006.04768*, 2020.
- [5] M. M. Rahman, M. Munir, and R. Marculescu, “EMCAD: Efficient multi-scale convolutional attention decoding for medical image segmentation,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [6] O. Ronneberger, P. Fischer, and T. Brox, “U-Net: Convolutional networks for biomedical image segmentation,” in *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*. Springer, 2015, pp. 234–241.
- [7] Z. Zhou, M. M. R. Siddiquee, N. Tajbakhsh, and J. Liang, “UNet++: A nested U-Net architecture for medical image segmentation,” in *Deep Learning in Medical Image Analysis (DLMIA/ML-CDS @ MICCAI)*, 2018, pp. 3–11.
- [8] O. Oktay *et al.*, “Attention U-Net: Learning where to look for the pancreas,” in *Medical Imaging with Deep Learning (MIDL)*, 2018.
- [9] F. I. Diakogiannis, F. Waldner, P. Caccetta, and C. Wu, “ResUNet-a: A deep learning framework for semantic segmentation of remotely sensed data,” *ISPRS Journal of Photogrammetry and Remote Sensing*, vol. 162, pp. 94–114, 2020.
- [10] N. Ibtehaz and M. S. Rahman, “MultiResUNet: Rethinking the U-Net architecture for multiresolution image segmentation,” *Neural Networks*, vol. 121, pp. 74–87, 2020.
- [11] J. Chen *et al.*, “TransUNet: Transformers make strong encoders for medical image segmentation,” *arXiv preprint arXiv:2102.04306*, 2021.
- [12] H. Cao *et al.*, “Swin-Unet: Unet-like pure transformer for medical image segmentation,” in *European Conference on Computer Vision (ECCV) Workshops*, 2022.
- [13] A. Hatamizadeh, Y. Tang, V. Nath, D. Yang, A. Myronenko, B. Landman, H. R. Roth, and D. Xu, “UNETR: Transformers for 3D medical image segmentation,” in *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2022, pp. 574–584.
- [14] H. Wang *et al.*, “UCTransNet: Rethinking the skip connections in U-Net from a channel-wise perspective with transformer,” in *AAAI Conference on Artificial Intelligence*, 2022, pp. 2441–2449.

- [15] R. Azad, M. T. Al-Antary, M. Heidari, and D. Merhof, "TransNorm: Transformer provides a strong spatial normalization mechanism for a deep segmentation model," *IEEE Access*, vol. 10, pp. 108 205–108 215, 2022.
- [16] R. Azad *et al.*, "DAE-Former: Dual attention-guided efficient transformer for medical image segmentation," in *Predictive Intelligence in Medicine (PRIME) Workshop @ MICCAI*. Springer, 2023, pp. 83–95.
- [17] S. Li, Q. Wang, G. Shi, K. Zhang, S. Wang, and Y. Zhang, "DBAANet: Dual-branch attention aggregation network for medical image segmentation," in *2025 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*. IEEE, 2025, pp. 5651–5658.
- [18] G. Ding, J. Wu, K. Yang, Z. Huang, and Z. Chen, "DAMamba: Dual-attention mamba for lightweight and robust breast ultrasound tumor segmentation," in *2025 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*. IEEE, 2025, pp. 5416–5423.
- [19] B. Landman *et al.*, "MICCAI multi-atlas labeling beyond the cranial vault workshop," in *MICCAI Workshop*, 2015.
- [20] D. Jha *et al.*, "Kvasir-SEG: A segmentation dataset and benchmark for gastrointestinal polyp segmentation," in *MultiMedia Modeling (MMM)*, 2020, pp. 451–462.
- [21] J. Bernal, F. J. Sánchez, G. Fernández-Esparrach, D. Gil, C. Rodríguez, and F. Vilariño, "WM-DOVA maps for accurate polyp highlighting in colonoscopy: Validation vs. saliency maps from physicians," *Computerized Medical Imaging and Graphics*, vol. 43, pp. 99–111, 2015.
- [22] N. Codella *et al.*, "Skin lesion analysis toward melanoma detection 2018: A challenge hosted by the international skin imaging collaboration," *arXiv preprint arXiv:1902.03368*, 2019.
- [23] H. Wang, S. Xie, L. Lin, Y. Iwamoto, X.-H. Han, Y.-W. Chen *et al.*, "Mixed transformer U-Net for medical image segmentation," in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2022, pp. 2390–2394.
- [24] O. Bernard, A. Lalande, C. Zotti, F. Cervenansky, X. Yang, P.-A. Heng, I. Cetin, K. Lekadir, O. Camara, M. A. G. Ballester *et al.*, "Deep learning techniques for automatic MRI cardiac multi-structures segmentation and diagnosis: Is the problem solved?" *IEEE Transactions on Medical Imaging*, vol. 37, no. 11, pp. 2514–2525, 2018.